Breaking substitution ciphers with Markov models

Algorithms description and full implementation in Python

Enrico Rollando*

Augustus 6, 2021

Abstract

Cryptography can be an interesting field to practice the knowledge of statistical methods. In this document I use the simple substitution cipher as toy-example to decrypt messages without knowing the key. I focus on a couple of unsupervised learning techniques based on Markov models, review the theory behind them and show a concrete code implementation and usage example of these algorithms.

Introduction

Can we break a simple substitution cipher automatically? "To break" here means: discover and use some weaknesses to restore - at least in part - the original text.
To define the scope, let's look only at plaintext and ciphertext composed in English capital letters and space, 26+1 symbols in total. Let's further say that the space is encrypted as space.

The *simple substitution cipher* is a mapping function $S_k(text) \rightarrow ciphertext$ where k is the key. It's a bijection between the alphabet sequence A and one of its |A|! permutations. Given the ciphertext only, I want to figure out the inverse mapping $S^{-1}(ciphertext)$ in an unsupervised way.

Let's try to define a hypothetical function: *fit(text)* that evaluates the current guess, ideally as "what is the probability that this text is the original plaintext?". A slightly different question that leads to a similar result and which I *can* actually define is: "how likely is it that this text is (English) plaintext?". I can use this function to evaluate and compare different guesses and choose the best one.

To model this evaluation I could use for example the unique features of a specific language. How often does a certain letter, or a certain group of letters (n-grams), appear in the plaintext. Or: how often a certain letter happens to follow another. And then look at the ciphertext and measure the distance using these statistics.

There is still a problem, the number of possible keys is huge (26! for the English alphabet), too many to try them all in a reasonable amount of time. Other techniques have to be used, in this document I show some of them.

Let us first look at the cipher to break and briefly present the Language Model which is the fundamental feature used to exploit the weakness of the cipher.

The simple substitution cipher

A simple substitution cipher [1] is a set of functions each modeling a one-to-one mapping between two alphabets: the plaintext- and the ciphertext- alphabet. This mapping is the key of the cipher.
Spaces are sometimes removed from plaintext before encryption, as they provide information about word boundaries, which in turn helps recovering the plaintext.

P=$[a_1, ..., a_{26}]$, example: ABCDEFGHIJKLMNOPQRSTUVWXYZ
C=$[c_1, ..., c_{26}]$, example: HQUVNCFYPXJDIMOBWLAZTGSRKE

---

*e.rollando@outlook.com

S: $f(P) \to C$ (1)

$S(\text{"THISISATEST"})=\text{"ZYPAPAHZNAZ"}$

$S^{-1}(\text{"ZYPAPAHZNAZ"})=\text{"THISISATEST"}$

The Language Model

In any language there are implicit (not formalized) rules that we use when we make or interpret a sentence. For example, an incomplete sentence "I like to write my letters by ..." would be expected to end with the word "hand". The initial words give a specific *context* and a context gives some words more *plausibility* than others.

At the character level, some knowledge of English is enough to read an incomplete word like "uncomf.rt.b.y". The plausibility can be interpreted statistically: what is the most likely character in this position, given the visible sequence of characters?

The Language Model is tied to a specific language and focuses on frequencies of elements, on character or word level. As an example I consider the bigram English Language Model LM_{EN}, informally:

- bigram frequency:
  - how frequent is a bigram ab in English text? $P_{LM}(ab) = \frac{count(\text{occourrences of ab})}{len(\text{text})-1}$
  - For a bigram ab how frequently does b follow a in English text? $P_{LM}(b|a) = \frac{count(ab)}{count(a)}$ (2)

The weakness

The simple substitution cipher encrypts the same symbol each time the same way. The ciphertext positionally maintains statistical information of the original plaintext. In other words, the i-th n-gram of both plaintext and ciphertext will have the same frequency in its message. For example, if the second bigram is a double-letter, the most probable plaintext bigram is "TT" or "LL".

To exploit this fact I could for example define a generative distribution G for the plaintext, controlled by some parameters Θ, then generate candidate plaintext, such that its positional n-grams distribution probability matches the one of the ciphertext.

This process could be progressively refined by automatically tweaking the parameters Θ to let distribution G converge to the observed distribution of the ciphertext.

Break the cipher

Let's review some theory first, which allows me to introduce an important amount of tools.

Break the cipher with the Expectation-Maximization (EM) algorithm

The main problem, at first sight, is the huge amount of keys to be tested to find a solution. This is a common problem for optimization algorithms and there are different approaches to solve it.

A well-known algorithm is the Expectation-Maximization (Dempster, Laird, Rubin, 1977) [2].

A very good description of the algorithm can be found in "Theory and Use of the EM Algorithm" (Maya R. Gupta and Yihua Chen, 2011) [3].

Some interesting research has been done in the last years about using EM in substitution ciphers, for a nice example read: "Unsupervised Analysis for Decipherment Problems" (Knight, Kevin & Nair, Anish & Rathod, Nishit & Yamada, Kenji, 2006) [4].

Expectation-maximization is, like *dynamic programming*, an optimization technique used to reduce the complexity of an algorithm by breaking it down into smaller and more tractable steps which are repeated until convergency is reached.

EM identifies a family of algorithms, in this document I use a specialized implementation of the EM for the discrete case.

EM basis: statistical foundation

Markov model A Markov model (Gagniuc, Paul A., 2017) [5] can be used for both discrete and continuous stochastic processes satisfying the Markov property: the previous state of the system is sufficient to describe its current state. This simplification allows for a series of techniques, such as Monte Carlo, which goal is to find out the parameters that stabilize the model without trying all the possible combinations.

An example of Markov model is the Markov chain. A Markov chain is a model that defines a conditional probability of moving to state x_i when the last k states were: $x_{i-k}...x_{i-1}$ and k is the level of "visibility" on the previous states.
For example, an LM can be approximated by a Markov chain of order n by using n-gram conditional probabilities (2).

The transition probability from the last sequence $p_{i-k}...p_{i-1}$ to the next letter p_i will be:

$$P(p_i|p_{i-k},...,p_{i-1})$$

The transition probability must define a valid probability distribution such that:

$$\sum_{p_i} P(p_i|p_{i-k},...,p_{i-1}) = 1$$

The Markov chain defines a process which states change over time, based on the last k-states and the transition matrix.

A Markov chain of order 1, in the discrete case, is defined by a set of states $(x_0,...x_n)$ and by a transition matrix:

$$A : p_{ij} = P(X_t = j|X_{t-1} = i)$$

which describes the probabilities to move from previous state i to current state j. This simplification minimizes the knownledge of the model by stating that the new transition only depends on the previous k states.

When not all the states of the process are observable, but the observations are a probabilistic function of the states, we talk about Hidden Markov models. This technique was defined by Leonard E. Baum and colleagues in the 1960's, a very good reference can be found in "A tutorial on hidden Markov models and selected applications in speech recognition" (L. R. Rabiner, 1989) [6].

We could model a substitution cipher with a HMM where the plaintext is the unknown state sequence, because it is hidden from the observation, and only the ciphertext is visible. The LM will help in evaluating the plausibility of the guessed plaintext as English text.

EM algorithm

The EM algorithm is a way to reduce the complexity of the problem of finding the best parameters of a model, maximizing the probability that the model would produce the observed output.

EM: A brief description There exist different forms of the algorithm, but a general explanation is the following:
The algorithm is used to improve the (possibly unknown or incomplete) parameters Θ of a given model by processing an unknown, not observable value X so that the output would maximally match some observed data Y.
It does this by repeating the following steps until the improvement is negligible (=the algorithm converged):

- E-step:
  - Using the current parameter values Θ, estimate X
- M-step:
  - using the estimated X, improve (find a better estimate for) Θ such that $P(Y|\Theta)$ grows.

For the cipher, X is the plaintext and Y is the ciphertext.

I'll reinstate the problem of solving the cipher, this time using the HMM.

EM: HMM model for the simple substitution cipher We have:

- a hidden sequence of plaintext characters $H : p_1, ..., p_n$
- an observed sequence of ciphertext characters $V : c_1 ... c_n$

In our case they use the same vocabulary of English letters.
The hidden sequence is drawn from the English LM statistics: it is English text and it is expected to respect the rules of the English language.
The observed sequence is bound to the hidden sequence by an as yet unknown transformation.
The goal is to find out the most likely sequence of plaintext characters that originated the ciphertext:

$$H = argmax_H P(H|V) = argmax_H(V, H)$$

from Bayes' theorem and simplifying, as the denominator has no part in the argmax.

The HMM is defined by:

- a fixed set of N hidden states X = {set of possible letters in plaintext}
- a fixed set of M observable states Z = {set of letters in the ciphertext}
- the probability Π of starting with a certain hidden state p_i
- the probability of moving from state p_i to p_j, kept in the so-called transition matrix A (size NxN).
- the probability of observing c_i being in state p_i, kept in the so-called emission matrix B (size NxM).

The hidden states are the unknown plaintext character sequence, the visible states are the observed ciphertext character sequence.
The parameters of the model are $\Theta = (\Pi, A, B)$.
These are the initialization steps:

- initialize Π from the LM with the probability of starting a sentence with letter i.
- initialize A with the probability of switching from letter i to letter j, taken from the bigram LM: $P(p_i|p_{i-1})$. This matrix will be kept constant during the optimization process as it's the middle used to exploit the cipher weakness.
- randomly initialize B which is the parameter to optimize. It represents the hidden distribution P(V|H) which "explains" how the plaintext characters get encrypted. When considering the space-plaintext symbol, set a one-to-one corrispondence with the space-output symbol giving it a probability of 1 in the matrix.

EM: the optimization process The EM algorithm receives as input the initial probability state distribution Π, the transition matrix A, the emission matrix B and the observations, then it tweaks the parameters to reach the maximum likelihood of the hidden variable H by alternating two steps:

1. E-step: using the evidence (the observed data) and the current parameters Θ infer the hidden values (those that the system will consider the most probable, given the current knowledge at this timestep).
2. M-step: using inferred hidden values re-estimate the parameters.
3. repeat from 1 until convergence - or until the desired maximum number of iterations is reached.

An important note is that the algorithm can converge to a local optimum. To improve the chance of finding the global optimum, the algorithm should be executed multiple times and the results compared.

The EM algorithm for a HMM can be implemented in different ways. Here I make use of a specialized implementation for discrete HMM: the Baum-Welch algorithm (from the inventors Leonard E. Baum and Lloyd R. Welch, 1960s), which in turn uses the forward-backward algorithm.

EM Implementation: the Baum-Welch algorithm Consider the hidden state H: $p_1 ... p_n$ which we just initially guess - by choosing B - and the observation V: $c_1 ... c_n$

- Each transition $p_{t-1} \rightarrow p_i$ happens with probability $P(p_t|p_{t-1})$. The c_i character is emitted with probability is $P(c_t|p_t)$.

- The probability of the visible sequence V given a particular hidden state sequence is the initial probability of starting with p_1, $\Pi(p_1)$ and then emitting c_1: $P(c_1|p_1)$, times the transition probabilities and the emission probabilities at each timestep t:
  $P(c_{1..N}, p_{1...N}|\Theta) = \Pi(p_1)P(c_1|p_1)\prod_t P(p_t|p_{t-1})P(c_t|p_t)$.
  Notice that this can be computed recursively.
- The probability of the visible sequence V is the marginalization of the previous computation for every possible hidden sequence $H = p_{1...N}$: $P(V|\Theta) = \sum_H \Pi(p_1)\prod_i P(p_t|p_{t-1})P(c_t|p_t)$. This is a computational challenge.
- Baum-Welch easies the computation by splitting it in two parts for each timestep, based on *what has already been observed* (forward pass) and *what it will be observed* (backward pass).
- The **forward pass** of the Forward-Backward algorithm reduces the complexity by computing $P(H_t = p_t, V_{1...t}|\Theta)$ recursively and then marginalizing the result. At each step t for each hidden state i and visible output $V_t = c_t$ the value *alpha* is defined as follows:
  $\alpha_t[i] = B[V_t]\sum_j \alpha_{t-1}(j)A[ij]$
- The **backward pass** will complete this partial result: where the forward pass computes the probability up to step t, the backward pass does the same for the subsequent timesteps, going backwards from step t to step $t+1$ and computing $P(V_{(t+1)...N}|p_t, \Theta)$. This pass can be again computed recursively in a similar way. For this computation the value *beta* is introduced:
  $\beta_t[i] = \sum_j \beta_{t+1}[j]B[j, V_{t+1}]A[ij]$
- The probability of moving from state p_t to p_{t+1} is proportional to $\alpha_t[i]A[ij]B[j, V_{t+1}]\beta_{t+1}(j)$. Splitting this computation with alpha and beta is the core optimization that avoids the algorithm being exponential.
- The conditional probability of the hidden states H is: $P(H|V, \Theta) = \frac{P(H,V|\Theta)}{P(V|\Theta)}$. The partial results of the forward and backward passes can be later reused to calculate this probability.
- Finally, Π, A and B can be updated by marginalizing $P(H|V, \Theta)$ and normalizing.

There are a couple of important details which I used in my implementation, (see: [6]: "Implementation issues for HMMs"):

- The elements of the transition matrix A should never be zero.
- The forward and backward algorithms are recursively computed products of probabilities. As the ciphertext length grows, the computation converges exponentially to 0, which generates underflows and makes the result unreliable. To avoid the problem, at each recursion step of alpha I scaled the numbers and I used the same factor to scale the correspondent step of beta.
- The system converges when there is no further improvement in $P(V|\Theta)$, which would be calculated directly from the alphas. But their values are normalized, so the scale factors can be used instead.
- I used $log(P(V|\Theta))$ to turn the product of very small numbers into sum of their logarithms.
- In recalculating B I forced the emission probability $space \leftrightarrow space$ to be 1.

EM implementation: a concrete example   So far I have summarized the algorithm and its Forward-Backward passes. Let's use it to automatically decrypt some text.

```
1  [import...]
2
3  class StringEncDec:
4      def ordToChar(n):
5          return ' ' if n==26 else chr(65+n)
6
7      def charToOrd(n):
8          return 26 if n==' ' else ord(n)-65
9
10 class BaumWelch:
11     def normalize(self,M):
12         tot=np.sum(M,axis=min(1,len(M.shape)-1),keepdims=1)
13         return M/tot,tot
14
```

```python
15      def alpha(self,A,B,PI,V,debug=True):
16          alphas,scale_facts=np.zeros([self.T,self.num_hidden]), np.zeros(self.T)
17          alphas[0],scale_facts[0]=self.normalize(PI*B[:,V[0]])
18          for t in range(1,self.T):
19              alphas[t],scale_facts[t]=self.normalize(B[:,V[t]]*
                    np.sum(alphas[t-1]*A.T,axis=1))
20          return alphas,scale_facts
21
22      def beta(self,A,B,V,scale_facts,debug=True):
23          betas=np.zeros([self.T,self.num_hidden])
24          betas[self.T-1]=np.array([1./scale_facts[self.T-1]])
25          for t in range(self.T-2,-1,-1):
26              betas[t]=np.sum(betas[t+1]*B[:,V[t+1]]*A,axis=1)/ scale_facts[t]
27          return betas
28
29      def process(self,A,B,PI,V,max_iter,recompute_A=True,
            recompute_B=True,recompute_PI=True,ext_fun=False,verbose_fun=False):
30          self.T,self.num_hidden,self.num_visible,log_scale_facts=V.size,
                B.shape[0],B.shape[1],None
31          for it in range(max_iter):
32              alphas,scale_facts=self.alpha(A,B,PI,V)
33              # finish if the prob of observations stops to increase
34              new_log_scale_facts=np.sum(np.log(scale_facts))
35              if new_log_scale_facts<(log_scale_facts or new_log_scale_facts):
36                  if verbose_fun: print("Algorithm has converged")
37                  break
38              log_scale_facts=new_log_scale_facts
39              betas=self.beta(A,B,V,scale_facts)
40              gammas=np.zeros([self.T,self.num_hidden, self.num_hidden])
41              for t in range(self.T-1):
42                  gammas[t]=(A.T*alphas[t,:]).T* (betas[t+1]*B[:,V[t+1]])
43              # sum all j's in gammas_i
44              gammas_i=np.sum(gammas,axis=2)
45              gammas_i[self.T-1]=alphas[self.T-1,:]
46              # optionally recompute A
47              if recompute_A:
48                  gammas_t=np.sum(gammas,axis=0)
49                  A=(gammas_t.T/np.sum(gammas_t,axis=1)).T
50              # optionally recompute B
51              if recompute_B:
52                  for v in range(self.num_visible):
53                      B[:,v]=np.sum(gammas_i[V==v]  ,axis=0)
54                  B=(B.T/np.sum(gammas_i,axis=0)).T
55              # optionally recompute PI
56              if recompute_PI:
57                  PI=gammas_i[0]
58              # optionally execute external function
59              if ext_fun:
60                  ext_fun(A,B,PI,V)
61              # optionally execute verbose function
62              if verbose_fun:
63                  verbose_fun(V,B,it,log_scale_facts)
64
65          return A,B,PI,log_scale_facts
```

EM usage example: computing statistics for the English bigrams LM   I used the Brown Corpus (Francis, W. Nelson & Henry Kucera, 1967) [7], upper-cased, from which I removed all non-alphabetic characters except space. Then I obtained the bigram statistics for the English LM using the Baum-Welch algorithm to compute the matrix A. Notice that it would be faster to compute this matrix directly from the text (2).

Parameters initialization   The algorithm must receive A: transition probability matrix; B: identity matrix; PI: a vector of initial state probabilities, V: the text to decrypt.
As the implementation expects a set of states $0 \ldots N$, I converted the characters $A, \ldots, Z$ to their ordinal numbers $0, \ldots, 25$ and the space to 26.
Keep B constant and get the bigram transition probabilities in A and the start probabilities in PI.
The entire corpus is slow to process this way, I limited the input to the first million characters and got acceptable results in speed and quality.

```python
1  def preprocess(text):
2      return re.sub('\s+',' ',re.sub('[^'+string.ascii_uppercase+']','
           ',text.upper())).strip()
3
4  # bigram Language Model using Baum-Welch
5  num_states=27
6  PI=np.random.rand(num_states)
7  # init A randomly
8  A=np.random.rand(num_states,num_states)
9  # B is identity
10 B=np.identity(num_states)
11
12 brownie=' '.join([preprocess(' '.join(sents)) for sents in brown.sents()])
13 V=np.array([StringEncDec.charToOrd(i) for i in brownie[:1000000]])
14 baum=BaumWelch()
15 A,_,PI,_=baum.process(A,B,PI,V,max_iter=150,recompute_B=False)
```

EM usage: decrypt a text   Now keep A and PI constant and get the emission probabilities in B.

```python
1  # a tweak for subst. cipher: force 1-1 mapping for space:space
2  def ext_fun(A,B,PI):
3      B[26,:]=0
4      B[:,26]=0
5      B[26,26]=1
6      B=baum.normalize(B)[0]
7
8  # show the current try: decode V using B
9  def verbose_fun(V,B,it,log_scale_facts):
10     print('#It',it,''.join([StringEncDec.ordToChar( np.argmax(B[:,v])) for v in
           V]),log_scale_facts)
11
12 # use Baum-Welch to decrypt this text
13 ciphertext='RBO RPKTIGO VCRB BWUCJA WJ KLOJ HCJD KM SKTPQO CQ RBWR LOKLGO VCGG CJQCQR KJ
       SKHCJA WGKJA WJD RPYCJA RK LTR RBCJAQ CJ CR '
14 V=np.array([StringEncDec.charToOrd(c) for c in list(ciphertext)])
15 # just a single run
16 for h in range(num_states):
17     for v in range(num_states):
18         B[h,v]=100 if v==26 and h==26 else random.random()
19 B=baum.normalize(B)[0]
```

```
20 _,BB,_,VAL=baum.process(A,B,PI,V,max_iter=5000,ext_fun=ext_fun,verbose_fun=verbose_fun,
       recompute_A=False,recompute_PI=False)
```

```
1 # output:
2 #It 0 BZN BQWNXLN JQBZ ZWVQYD WY WGNY JQYY WE PWNQDN QD BZWB GNWGLN JQLL QYDQDB WY
       PWJQYD WLWYD WYY BQZQYD BW GNB BZQYDD QY QB  -374.15396169450446
3 #It 1 BHN BQPNXLN JABH HAZARD AR PCNR JARY PE PPNQDN AD BHAB CNPCLN JALL ARDADB PR
       PPJARD ALPRD ARY BQZARD BP CNB BHARDD AR AB  -312.11160093692797
4 #It 2 THF TOPXXLF JATH HAZARD AR PCFR MARY PE WPXONF AN THAT CFPCLF JALL ARNANT PR
       WPMARD ALPRD ARY TOZARD TP CXT THARDN AR AT  -297.80567428378487
5 [...cut...]
6 #It 15 THE TOXUPKE WITH HAZIND AN XQEN MING XF JXUOSE IS THAT QEXQKE WIKK INSIST XN
       JXMIND AKXND ANG TOVIND TX QUT THINDS IN IT  -243.48633910060906
7 #It 16 THE TOXUPLE WITH HAZIND AN XQEN MING XF JXUOSE IS THAT QEXQLE WILL INSIST XN
       JXMIND ALXND ANG TOVIND TX QUT THINDS IN IT  -242.68538455928874
8 [...cut...]
9 #It 950 THE TROUPLE WITH HAZING AN OQEN MIND OF JOURKE IK THAT QEOQLE WILL INKIKT ON
       JOMING ALONG AND TREING TO QUT THINGK IN IT  -237.72211809668678
10 Algorithm has converged
```

This run quickly converged to a fairly understandable solution. There are still a few mistakes in the mapping, but the plaintext can be easily guessed now.

Unfortunately, however, the algorithm generally falls into local optima. In that case, just repeat the run a few times by reinitializing B on each run and using VAL $\leftarrow log(P(V|\theta))$ to choose the best results.

Decoding the results   In decoding the results, I directly used the matrix B to find the index of the best emission state for a given symbol v: argmax(B[:,v]). This was enough for the purposes. To answer the (different) question of "what is the most likely hidden sequence state for the given emission matrix", the Viterbi algorithm should be used (A.Viterbi, 1967) [8].

EM: Review of the Baum-Welch approach   With the Baum-Welch algorithm I can automatically decrypt a text in seconds.
The algorithm is very flexible as any of its parameters A, B, Θ can be optimized. I used the algorithm to prepare the English bigram LM and used it again to decrypt some text.

let's make some considerations:

- the entropy of an English text tells us how much information is given us by a text. This can be used to compute the *unicity distance*: (Claude Shannon, "Communication Theory of Secrecy Systems", 1949) [9]: the minimum length of text which gives us enough information to decrypt, which is defined as the entropy of the keyspace divided by the per-character redundancy in bits (ca. 3.2 for English text):

$$U = \frac{H(K)}{D} = \frac{log2(26!)}{3.2} \approx 28$$

  For the Baum-Welch approach to work, much longer text is needed.
- this technique only exploits the bigram LM, or first-order HMM, which gives to the statistics a fair restricted horizon visibility.
- it's possible to improve the results by modifying the algorithm and make use of higher-order HMMs. Various solutions can be further used to improve speed and memory usage, but the algorithm will be slower and memory intensive when using higher-order HMMs.
- the substitution cipher realizes a bijectional mapping, but I could not directly enforce nor exploit this knowledge in the algorithm.
- it's possible to give a hint to the process: like for the space-to-space mapping, the A matrix (randomly filled in the example) can hold higher probabilities for known or expected mappings. However, there

is no immediate way to tell the process that there are high expectations for specific words of part of sentences, like greetings or names.

Break the cipher with Metropolis-Hastings

Let's try out other solutions for the problem.
Is there a way to enforce the bijectional property of the cipher while solving the problem?
I briefly describe the Monte Carlo technique and the Metropolis-Hastings algorithm before using it with a new model.

Monte Carlo

When we have a distribution which we can estimate, or from which can be sampled, if we want to find related quantities which for some reasons are intractable or not easy to compute, we can spare the calculations and just approximate these quantities by random sampling.

As a typical example, we could estimate the value of π by just drawing a circle of radius 1 inscribed in a 2x2 square and then repeatedly generate a pseudo-random point within the square area (coordinates $x, y \in [1, 1]$) and measure the proportion of times the point falls within the circle area.

Knowing the relation between the area of a square and the area of the inscribed circle:

$$\text{square area} = 2r * 2r = 4$$

$$\text{circle area} = \pi * r^2 = \pi$$

$$\frac{\text{circle area}}{\text{square area}} = \pi/4$$

We reach an estimation of the expected value of $\pi/4$ and, most importantly, for the law of large numbers this value will converge to the expected value $\pi/4$ for large n.

```
1 limit=1000000 # the more the better
2 print(sum([random.uniform(-1,1)**2 + random.uniform(-1,1)**2<=1 for i in range(limit)
      ])*4.0/limit)
3 3.14188
```

Under the "Monte Carlo" methods fall different techniques, which stochastically tweak the parameters of the model to explore unknown values and evaluate its outputs, such that the model can be improved. In a HMM this finetuning is called "random walk".
This walk is mostly done by generating random samples and discarding not plausible ones (rejection sampling) or by generating samples from some weighted distribution (importance sampling).
For example, in the case of a substitution cipher, consider a random walk where at each step two letters of the cipher key are swapped and the resulting plaintext is sampled to evaluate its plausibility.

Metropolis-Hastings, brief introduction

Metropolis-Hastings is a Monte Carlo rejection sampling algorithm that works well for the substitution cipher case.

- consider an unknown distribution P(x) which we however can estimate using only the observations.
- we have a function Pl(x) which is just proportional to the density of P(x).
- we make a random walk to progressively improve the unknown function by randomly make a change on the parameters of the model and then sample from the results.
- the new value will be accepted or refused using the "acceptance function" which compares Pl(x) with its previous value to decide.

In the longer term (and under some basic conditions) this walk is guaranteed to optimize the system and return samples which follow the unknown distribution P(x).

More about convergency of Markov models

A (irreducible, aperiodic) discrete Markov model with states distribution $\Pi(x)$ and transition distribution $P(y|x)$ will converge when:

$$\sum_x \Pi^*(x)P(y|x) = \Pi^*(y)$$

That is, the probability of being in y is equal to the probability of getting into y from any state x.
Π^* is the stationary distribution here (it may not exist or may not be unique).
Starting with the initial state distribution Π and transitioning n times using the transition matrix A, the model will enter a new state: $\Pi * A^n$.

For high values of n, Π becomes insignificant and the result will converge:

$$\Pi^* = \lim_{n\to\infty} A^n \Rightarrow \Pi^* * A = \Pi^*$$

That is, if the probability Π has not changed after the transformation, the chain has converged (the random walk has ended, the system has not changed state).

A theorem says that if $A_{xy} > 0 \quad \forall x, y \longrightarrow \exists$ unique Π^* (=if the transition matrix contains no zeroes, a unique stationary distribution exists).

The stationary distribution Π^* is equivalent to the unknown P(x) and it's the desired result of the optimization process.
A useful way to see that the stationary distribution exists is when the system is time-reversible, meaning (*detailed balance equation*):

$$\Pi(x)P(y|x) = \Pi(y)P(x|y)$$

$$(3)$$

In fact, summing for x at both sides:

$$\sum_x \Pi(x)P(y|x) = \sum_x \Pi(y)P(x|y) = \Pi(y)\sum_x P(x|y) = \Pi(y)$$

Some clues behind the inner workings of Metropolis-Hastings for the substitution cipher

So far, the HMM has represented the probability of each individual letter in plain text to be transformed and observed. The transformation was not bijective and the first-order LM was a limitation.
Let's rethink the model in a way that enforces the substitution mapping and simplifies the use of higher-order LM.

The technique discussed here comes from the very interesting "The Markov Chain Monte Carlo Revolution" (Diaconis, Persi, 2009) [10] which provides deeper insights into MCMC and uses Metropolis-Hastings to, among other things, reveal the contents of an encrypted message from a prison inmate.

Consider (1) the substitution cipher mapping functions $S : f(P) = C$ and the inverse mapping $S^{-1} : f^{-1}(C) = P$ and its correspondence to the distribution functions P(x|y) and P(y|x) where X is the space of the possible symbols and $x, y \in X$
Then the stationary distribution Π would represent the probability that the mapping function is the correct one, as applying successive transformations will not change nor improve the model and the model will output a sample of this stationary distribution, which is the mapping of the cipher.

The Π distribution is initially unknown, but we can easily obtain a probability distribution which is proportional to this distribution.
In fact, using the statistics from the n-gram LM, to estimate the plausibility of a sentence as English text we can compute the joint probability P_{LM} of the n-grams in the text sequence $\{c_1, ..., c_n\}$ transformed into candidate plaintext by the substitution S^{-1}:
(Plausibility:) $Pl = \prod_{i=1}^{n} P_{LM}(S^{-1}(c_{i+1})|S^{-1}(c_i))$

Metropolis-Hastings (symmetric proposal): the algorithm

The original Metropolis algorithm used the "symmetric proposal", which derives directly from the detailed balance equation.

It needs a symmetric function to choose a new candidate and another function to tune the results (some fitness function which in turn uses information from the LM) and it works as follows:

    0. start with
- some current state x (may be randomly generated)
- an symmetrical distribution: g(x|y)=g(y|x)
- a function Pl(x) (the plausibility), proportional to the desired target $\Pi(x)$

    1. from current x draw a new state y from the distribution g(y|x)
    2. calculate the acceptance rate $\alpha(y, x) = min\left(1, \frac{Pl(y)}{Pl(x)}\right)$
    3. Accept/Reject the candidate:
- accept the candidate and move to y if $\alpha >= \mu(0, 1)$ (not lower than a random uniform between 0 and 1)
- refuse otherwise and stay in x

    4. repeat from 1 until convergency

The acceptance ratio comes from (3):

$$\frac{P(y|x)}{P(x|y)} = \frac{\Pi(y)}{\Pi(x)}$$

using g to generate samples and alpha to accept or refuse the swap, rewrite the equation as:

$$\frac{g(x|y)\alpha(x, y)}{g(y|x)\alpha(y, x)} = \frac{\Pi(y)}{\Pi(x)}$$

because g is symmetric this simplifies to:

$$\frac{\alpha(x, y)}{\alpha(y, x)} = \frac{\Pi(y)}{\Pi(x)}$$

and because $Pl(x)$ is by requirement proportional to $\Pi(x)$ the condition can be fullfilled by choosing alpha to be:

$$\alpha(x, y) = min\left(1, \frac{Pl(y)}{Pl(x)}\right) = min\left(1, \frac{\Pi(y)}{\Pi(x)}\right)$$

Metropolis-Hastings implementation: a concrete example

With this algorithm, I am exploiting an LM with multiple n-grams. This is now easy due to the mild requirements for the Plausibility function. I don't even have to calculate a probability here. Consider that for each n-gram in the candidate text, the frequency of this n-gram in the plaintext will be higher if the n-gram has higher probabilities for that language. The sum of weighted n-gram frequencies (where weight is proportional to n-gram length as statistics on longer n-grams give us better information) is sufficient. Actually, "Frequency" is a normalized quantity, I don't even need to normalize here as long as the normalization factor (the sum of n-gram counts) stays the same throughout the process, because at each timestep the normalization factor cancels out when the plausibilities Pl_t/Pl_{t+1} are compared.

Furthermore, at each step an existing mapping will be swapped, this maintains the bijection between Plaintext and Ciphertext alphabet.

The n-gram matrix representing the LM is sparse for larger n-gram lengths and it would be inefficient to simply store it in an array. In this python example I am using a *dictionary*.

Finally, the algorithm's convergence technique will allow it to jump from one local minimum to another. For this reason, there is no explicit way to determine whether it has completed. For my fairly simple purpose of breaking the cipher, I just set a maximum number of iterations; another possibility is to stop the algorithm if there is no improvement for a number of iterations.

Here is my full implementation of the Metropolis algorithm and a concrete example:

```python
[import...]

def count_ngrams(text,min_dgram_len=2,max_dgram_len=9):
    return {dlen:
                {k:math.log(v) for k,v in Counter(text[idx : idx + dlen] for idx in
                    range(len(text)-dlen+1)).items()}
                for dlen in range(min_dgram_len,max_dgram_len+1)}

brownie=' '.join([preprocess(' '.join(sents)) for sents in brown.sents()])
mm=count_ngrams(brownie)

def metropolis(ciphertext):
    def plausibility(text,mm):
        cnt=count_ngrams(text) # count the ngrams in text, return logs of counts for
            ngrams of different length.
        return sum([mm[ngram_len][k]*ngram_len for ngram_len in cnt.keys() for k,v in
            cnt[ngram_len].items() if k in mm[ngram_len]])

    best=pl_f=plausibility(ciphertext,mm)
    fixed=set(ciphertext)-set(string.ascii_uppercase) # these symbols will not be
        transposed
    used_symbols=list(set(ciphertext)-fixed)
    all_symbols=set(string.ascii_uppercase) # only transposable chars
    smallprob=math.log(1e-3)
    for i in range(int(1e6)):
        c1=random.choice(used_symbols) # choose only between symbols actually used in
            the ciphertext
        c2=random.choice([c for c in all_symbols if c!=c1]) # switch off with this
        candidate=ciphertext.translate(str.maketrans({c1:c2, c2:c1}))
        pl_f_star=plausibility(candidate,mm) # evaluate the candidate
        mu=math.log(random.uniform(0,1))
        if mu<max(min(pl_f_star-pl_f,0),smallprob):
            if pl_f_star>best:
                best=pl_f_star
                print('It#',i, candidate)
            ciphertext=candidate
            pl_f=pl_f_star
            used_symbols=list(set(ciphertext)-fixed)

metropolis(' K CXJO XI HBD BKIO XL ENJHB HEN XI HBD CSLB ') # break a challenging short
    ciphertext
```

```
It# 1   K CXJO XI HBD BKIO XS ENJHB HEN XI HBD CLSB
It# 2   K CXJO XI TBD BKIO XS ENJTB TEN XI TBD CLSB
It# 3   K CXJO XI TBR BKIO XS ENJTB TEN XI TBR CLSB
It# 4   K CVJO VI TBR BKIO VS ENJTB TEN VI TBR CLSB
It# 7   M CVJO VI TBR BMIO VS ENJTB TEN VI TBR CLSB
It# 9   M CFJO FI TBR BMIO FS ENJTB TEN FI TBR CLSB
It# 15  M CFJS FI TBR BMIS FO ENJTB TEN FI TBR CLOB
It# 19  M COJS OI TBR BMIS OF ENJTB TEN OI TBR CLFB
It# 20  L COJS OI TBR BLIS OF ENJTB TEN OI TBR CMFB
It# 24  L COBS OF TJR JLFS OI ENBTJ TEN OF TJR CMIJ
It# 26  L COBS OF TKR KLFS OI ENBTK TEN OF TKR CMIK
```

```
12 It# 36   L CORS OF TKB KLFS OI ENRTK TEN OF TKB CMIK
13 It# 47   L CORS OF TIB ILFS OK ENRTI TEN OF TIB CMKI
14 It# 58   L CONS OF TIB ILFS OK ERNTI TER OF TIB CMKI
15 It# 70   L CONS OF TIK ILFS OB ERNTI TER OF TIK CMBI
16 It# 72   L CONS OF TIK ILFS OB ERNTI TER OF TIK CHBI
17 It# 113  L CONS OF THK HLFS OB ERNTH TER OF THK CIBH
18 It# 116  Y CONS OF THK HYFS OB ERNTH TER OF THK CIBH
19 It# 121  Y IONS OF THK HYFS OB ERNTH TER OF THK ICBH
20 It# 122  Y IONS OF THE HYFS OB KRNTH TKR OF THE ICBH
21 It# 124  Y IONS OF THE HYFS OR KBNTH TKB OF THE ICRH
22 It# 127  Y IONS OF THE HYFS OR KWNTH TKW OF THE ICRH
23 It# 137  Y IONS OF THE HYFS OR AWNTH TAW OF THE ICRH
24 It# 138  Y IONS OF THE HYFS OR ADNTH TAD OF THE ICRH
25 It# 146  Y IONS OF THE HYFS OR ADNTH TAD OF THE ILRH
26 It# 192  Y IODS OF THE HYFS OR ANDTH TAN OF THE ILRH
27 It# 262  P MOYS OF THE HPFS OR ANYTH TAN OF THE MIRH
28 It# 332  U MOYS OF THE HUFS OR ANYTH TAN OF THE MIRH
29 It# 789  I PONS OF THE HIFS OR ALNTH TAL OF THE PURH
30 It# 1175  A CONS OF THE HAFS OR IMNTH TIM OF THE CURH
31 It# 1878  A GIRD IN THE HAND IL WORTH TWO IN THE GULH
32 It# 1904  A CIRD IN THE HAND IL WORTH TWO IN THE CULH
33 It# 1934  A CIRD IN THE HAND IF WORTH TWO IN THE CUFH
34 It# 1958  A BIRD IN THE HAND IF WORTH TWO IN THE BUFH
35 It# 2116  A BIRD IN THE HAND IS WORTH TWO IN THE BUSH
```

Patristocrats

To make the cipher more difficult to break, at the expense of possible ambiguities in the interpretation of the plaintext, all spaces may be removed from the text. The American Cryptogram Association calls this kind of substitution code "Patristocrat".

To solve a Patristocrat, we need an LM calculated from a corpus *without* spaces. The same algorithms as before can be used, but longer ciphertext is needed to provide sufficient statistical input.

The higher-order HMM implementations, such as the one used for Metropolis-Hastings, give here much better results than the first-order implementations, as expected.

After the plaintext has been restored, it is also possible to restore the spaces - this time using an LM from a corpus *with* spaces. There are of course several ways to achieve the goal, such as dynamic programming, but - with only minor changes from the previous implementation - here I use Metropolis for this:

```python
def metropolis(ciphertext):
    def plausibility(text,mm):
        cnt=count_ngrams(text) # count the ngrams in text, return logs of counts for
            ngrams of different length.
        return sum([mm[ngram_len][k]*ngram_len for ngram_len in cnt.keys() for k,v in
            cnt[ngram_len].items() if k in mm[ngram_len]])

    best=pl_f=plausibility(ciphertext,mm)/len(ciphertext) # evaluate the current text
    for i in range(5000):
        c=random.choice(range(1,len(ciphertext)-2))
        candidate=list(ciphertext)
        if ciphertext[c]==' ':
            candidate=candidate[:c]+candidate[c+1:]
            pl_f_star=plausibility(''.join(candidate),mm)/ len(candidate) # evaluate the
                candidate with a space less
        elif ciphertext[c-1]!=' ' and ciphertext[c+1]!=' ':
```

```python
14              candidate[c:c]=' '
15              pl_f_star=plausibility(''.join(candidate),mm)/ len(candidate) # evaluate the
                    candidate with space more
16          mu=math.log(random.uniform(0,1))
17          if mu<min(pl_f_star-pl_f,0):
18              ciphertext=''.join(candidate)
19              pl_f=pl_f_star
20              if best<pl_f_star:
21                  best=pl_f_star
22                  ret=ciphertext
23                  print(ciphertext)
24      return ciphertext,best
25
26 metropolis(' IFYOULOOKFORPERFECTIONYOULLNEVERBECONTENT ') # restore spaces in this text
```

```
1 # output:
2  IFYOU LOOKF ORPERFECTIONYOULLNEVERBECONTENT
3  IFYOU LOOKF ORPERFECTIONYOULLNEVERBE CONTENT
4  IFYOU LOOKF OR PERFECTIONYOULLNEVERBE CONTENT
5  IF YOU LOOKF OR PERFECTIONYOULLNEVERBE CONTENT
6  IF YOU LOOKF OR PERFECTION YOULLNEVERBE CONTENT
7  IF YOU LOOKF OR PERFECTION YOULL NEVERBE CONTENT
8  IF YOU LOOKFOR PERFECTION YOULL NEVERBE CONTENT
9  IF YOU LOOK FOR PERFECTION YOULL NEVERBE CONTENT
10  IF YOU LOOK FOR PERFECTION YOU LL NEVERBE CONTENT
11  IF YOU LOOK FOR PERFECTION YOU LL NEVER BE CONTENT
12
13 (' IF YOU LOOK FOR PERFECTION YOU LL NEVER BE CONTENT ', 166.04992547413215)
```

Review of the Metropolis approach

Given enough time, very short text can be successfully decrypted with Metropolis.
The algorithm is quite flexible and allows the use of LM from higher order Markov models with ease.
The convergence rate is determined by the function that selects a new candidate and acceptance function
(which determines the rejection rate), modified versions of these functions form the basis for alternate versions
of this algorithm.

References

[1] Substitution cipher. (2021, July 27). In Wikipedia.
https://en.wikipedia.org/w/index.php?title=Substitution_cipher&oldid=1035776944

[2] Dempster, A.P.; Laird, N.M.; Rubin, D.B. (1977). "Maximum Likelihood from Incomplete Data via the EM
Algorithm". Journal of the Royal Statistical Society, Series B. 39 (1): 1–38. JSTOR 2984875. MR 0501537.

[3] Maya R. Gupta and Yihua Chen (2011), "Theory and Use of the EM Algorithm", Foundations and Trends®
in Signal Processing: Vol. 4: No. 3, pp 223-296. https://dx.doi.org/10.1561/2000000034

[4] Knight, Kevin & Nair, Anish & Rathod, Nishit & Yamada, Kenji. (2006). Unsupervised Analysis for
Decipherment Problems. https://dx.doi.org/10.3115/1273073.1273138.

[5] Gagniuc, Paul A. (2017). Markov Chains: From Theory to Implementation and Experimentation. USA,
NJ: John Wiley & Sons. pp. 1–256. ISBN 978-1-119-38755-8

[6] L. R. Rabiner, "A tutorial on hidden Markov models and selected applications in speech recognition," in
Proceedings of the IEEE, vol. 77, no. 2, pp. 257-286, Feb. 1989. https://dx.doi.org/10.1109/5.18626.

[7] Francis, W. Nelson & Henry Kucera. 1967. Computational Analysis of Present-Day American English. Providence, RI: Brown University Press.

[8] A. Viterbi, "Error bounds for convolutional codes and an asymptotically optimum decoding algorithm," in IEEE Transactions on Information Theory, vol. 13, no. 2, pp. 260-269, April 1967, doi: 10.1109/TIT.1967.1054010.

[9] Shannon, Claude. "Communication Theory of Secrecy Systems", Bell System Technical Journal, vol. 28(4), page 656–715, 1949. https://doi.org/10.1002/j.1538-7305.1949.tb00928.x

[10] Diaconis, Persi. (2009). The Markov Chain Monte Carlo Revolution. Bulletin of the American Mathematical Society. 46. 179textendash205. https://dx.doi.org/10.1090/S0273-0979-08-01238-X.